


Direct LLM API: GPT-4.1 E9 dry3 iterations

T=0 | max_output_tokens=8192 | timeout=600 s | workers=1 | max_retries=5 | retry_backoff=2 s

Step Exact Step+Module All Correct --- Best Step in this series



One case changes accuracy by 33.3 percentage points. This is a development screen, not GAIA-50.
All metric denominators=3. Failed outputs remain in the denominator; frozen upstream is marked *.